



Everybody be Cool, This is a Robbery!

**Jean-Baptiste Bedrune**, Gabriel Campana

firstname.lastname@ledger.fr

Hong Kong - New York - Paris - San Francisco - Vierzon

## Disclaimer

---


- The Donjon (Ledger Security Team) assess the security of technologies used by Ledger
- The vulnerabilities in this presentation were found during a security audit
- We don't want to single out one particular vendor
- Goals:
  - Raise awareness about HSM security
  - Lay the groundwork for other security researchers
  - Improve the overall security

# Agenda

---

- What is an HSM?
- Characteristics of the HSM assessed
- Brief intro to PKCS #11
- Developing tools for vulnerability discovery
- Vulnerability research and exploitation

HSM?



## What is an HSM?

---

- Security enclaves to store and process sensitive data
  - Computes cryptographic operations
  - Generate keys
  - Keys never leave the enclave
- Physical computing device:
  - PCI card or network appliance
  - One or more crypto-processors
  - Anti-tampering countermeasures

## Usage Examples

---

- PKI:
  - CA's private key generation and storage, certificates signing
  - Requirement for all CAs (CA-Browser Forum Baseline)
- Banking: CVV verification, transaction authorization, payment card personalization
- Telecommunications: strong cryptographic material for key injection by SIM manufacturer
- DNSSEC: storage of Root Zone keys (FIPS 140-2 level 4 HSM)
- Cloud services: encryption/decryption of customer data
- HSM-as-a-Service: Google, Microsoft, Amazon, etc.

## How much does it cost?

---

- Only a few vendors, no market share information
- No public prices, large range of models for each vendor
- According to [Hackable Security Modules](#) (REcon Brussels 2017):
  - Brand X, Model A: \$29,500.00
  - Brand Y, Model B: \$9,500.00
  - Brand Z, Model C: \$15,000.00


# Appliance (Host + HSM)

## HSM

- PCI Express card
  - Also available as a network appliance)
- FIPS 140-2 level 3 certified
- Components are coated in epoxy
- USB and serial ports for an optional smart-card reader
- Ethernet controller without connector

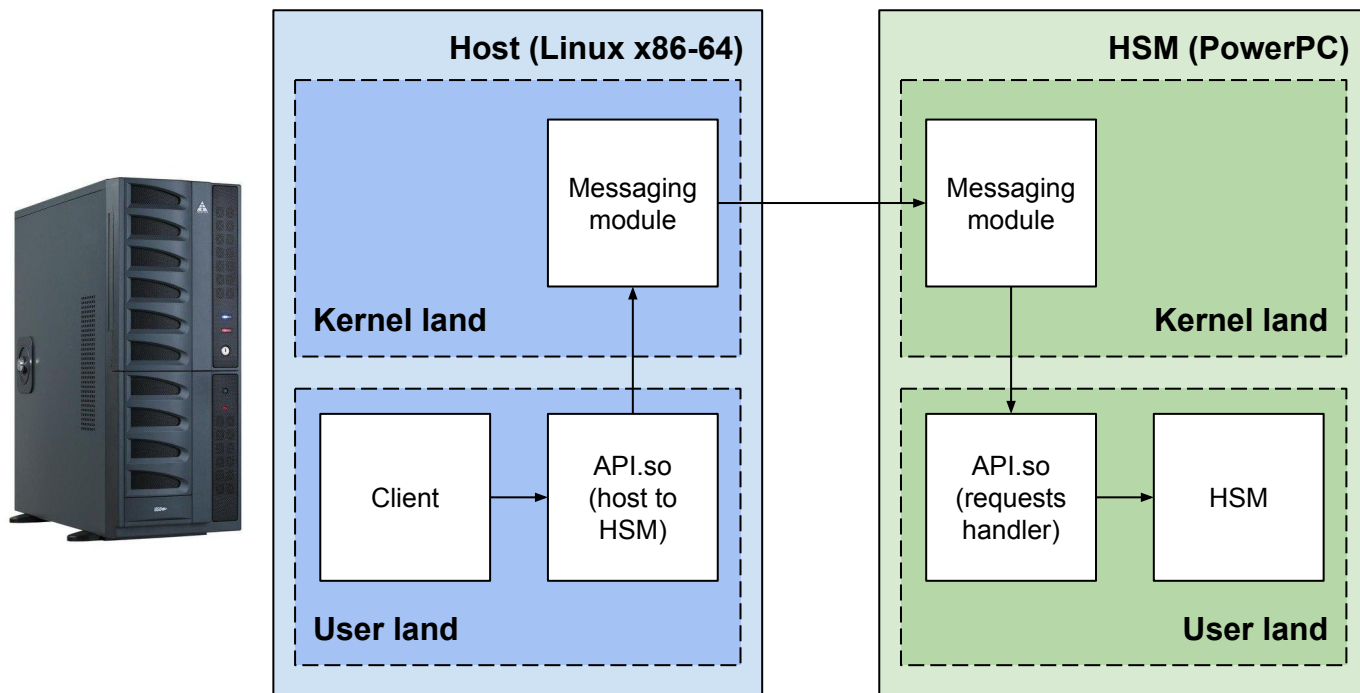
## Host

- Standard Linux server
- Linux Kernel modules
- CLI and GUI software, SDK





# Communication: Shared DRAM




## FIPS 140-2: Security Requirements for Cryptographic Modules

---

- U.S. government computer security standard
- Level 1: basic software requirements
- Level 2, 3, 4: physical requirements
  - Level 3: Detection and response to attempts at physical access, use or modification of the cryptographic module

Not a certification about software attacks

PKCS #11

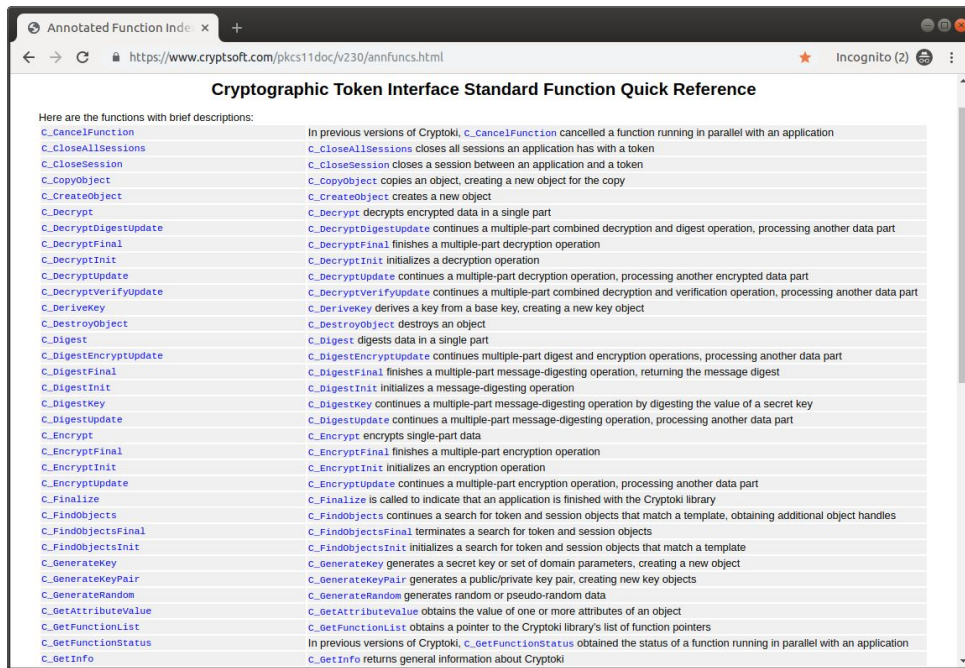


## PKCS #11: Introduction

---

- Generic interface to communicate with a cryptographic device
  - Smart card
  - HSM, etc.
- Portable API: Cryptographic Token Interface (Cryptoki)
  - Session management
  - Cryptographic objects manipulation
  - Operations on these objects (encryption, decryption, signature, etc.)

# Cryptographic Token Interface



- Few exposed functions (~70)
- But > 300 standard mechanisms, + proprietary mechanisms.

# Mechanisms

---

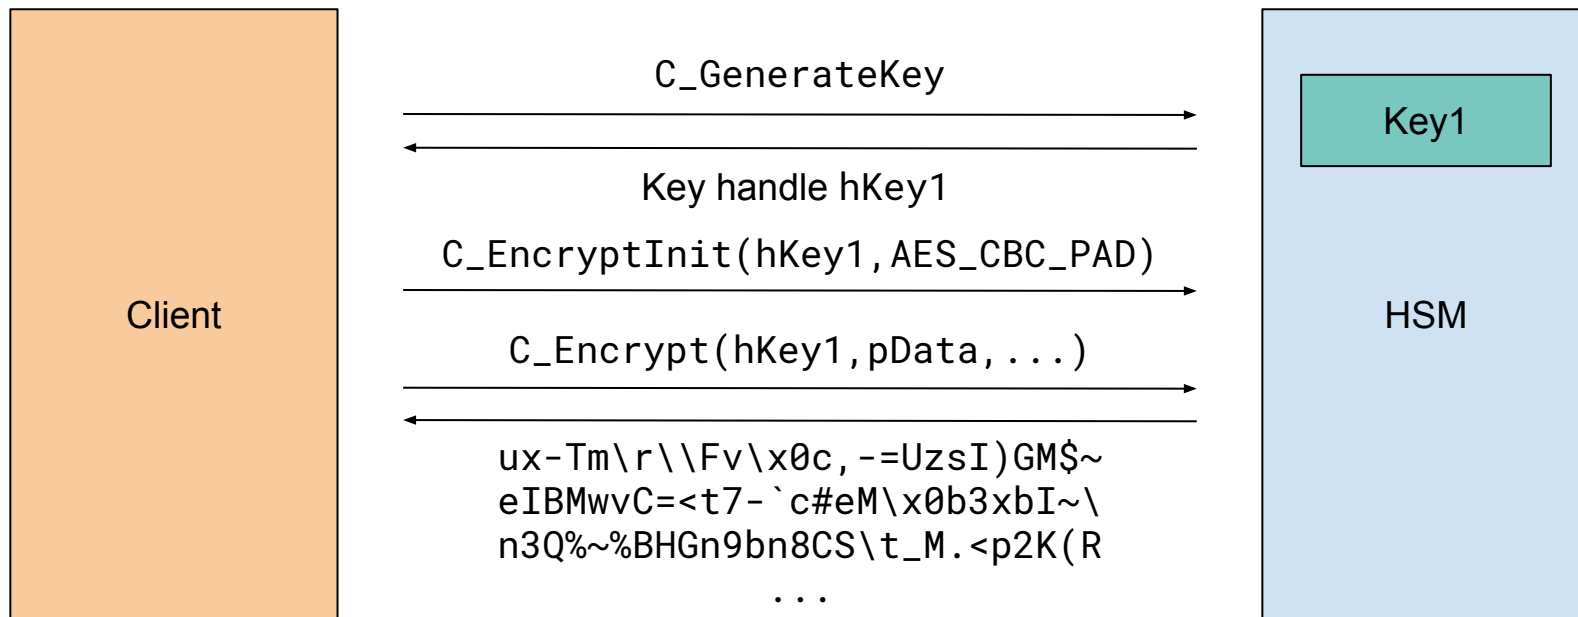
- Define how to perform a cryptographic operation.
- Mechanisms for encryption, decryption, hashing, wrapping, etc.
- Depends on the device. HSM: many mechanisms.
  - CKM\_SHA512\_HMAC
  - CKM\_RSA\_PKCS\_KEY\_PAIR\_GEN
  - CKM\_WRAPKEY\_AES\_CBC
  - CKM\_AES\_GCM
  - Mechanisms for telecom, banking...
- Some mechanisms take parameters: IV, salt, etc.

# Objects

---

- 3 types
  - Keys: secret, public, private
  - Certificates
  - Data: DSA / ECDSA parameters, etc.
- Cryptoki manipulates **objects** through their **handles**

# Objects




Key values are (usually) not sent to the host




# Object attributes

---

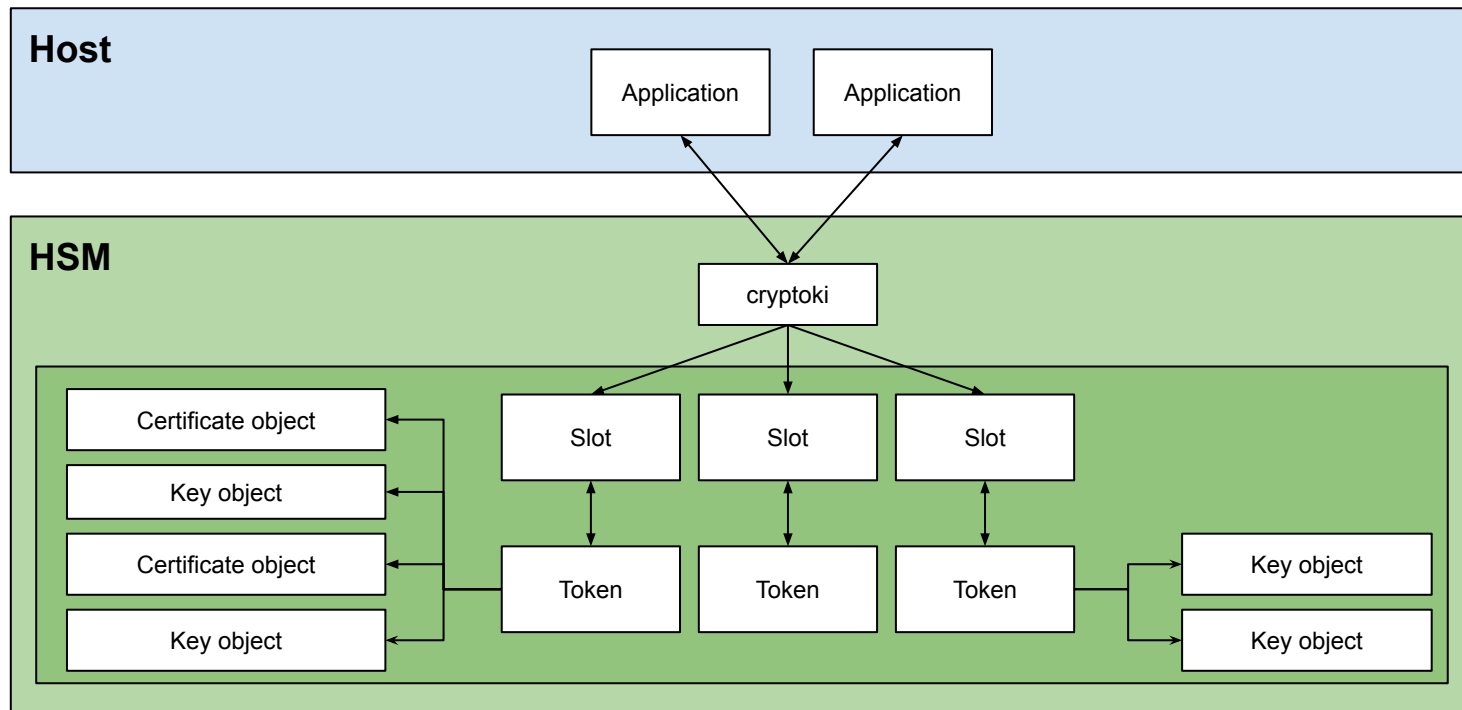


# Attributes for Security

- **Sensitive:** value cannot be extracted in plain text, must be wrapped
- **Not extractable:** value cannot be exported
- **Private:** user has to be logged to access the object



# Slots and tokens



# Threats

---

- Unauthenticated attackers gain access to private objects
- Attackers extracts keys marked as non-extractable
- Authenticated attackers gain access to other slots
- ...

Our goal: access to all objects from all slots, without authentication

## State of the art

---

- On the security of PKCS #11 (Clulow, CHES, 2003)
  - Much information on PKCS #11 security model
  - Encrypt then wrap, weak mechanisms...
- Your Bitcoin Wallet may be at risk: SafeNet HSM key-extraction vulnerability (Cem Paya, Gemini, 2015)
  - Weak mechanism in PKCS #11, enabled by default by SafeNet
- Hackable Security Modules - Reversing and exploiting a FIPS 140-2 Level 3 HSM firmware (Fotis Loukos, REcon Brussels, 2017)
  - Exotic CPU, focused on reverse engineering

Our contribution: attacks on PKCS #11 implementations

# Vulnerability Research and Exploitation

---

# Threat Model

---

- The attacker is able to execute commands on the host
  - Insider threats
  - Malicious data center employee with physical access
  - Administrator account compromise
  - Software vulnerabilities on the host
  - etc.

## Firmware

---

- CD-ROM provided by the vendor
  - CLI and GUI software
  - PKCS #11 API examples in C
  - Documentation for developers and administrators
  - Firmware update: signed, unencrypted, Linux 2.6.28.8 for PowerPC (2009)
- 
- Big Cryptoki library, very few other files
  - Few weeks of reverse engineering



# Tooling

---

# Module

---

- Unexpected option: new features can be added thanks to custom *modules*
- Expected usage:
  - PKCS #11 functions hook
  - New handlers on custom messages
- Not a vulnerability: requires admin privileges for loading
- Internals:
  - The SDK along a toolchain produces PowerPC ELF binaries from C source code
  - Modules loaded into the main process thanks to `dlopen()`
  - No `libc`

# Shell

---

```
user@host:~$ ./module-shell --init
```

```
[*] uploading busybox-powerpc to /sbin/busybox
```

```
[*] creating symlinks (might take a few seconds)
```

```
user@host:~$ ./module-shell id
```

```
uid=0 gid=0
```

```
user@host:~$ ./module-shell ps fauxwww
```

PID	USER	TIME	COMMAND
-----	------	------	---------

1	0	0:00	/init
---	---	------	-------

2	0	0:00	[kthreadd]
---	---	------	------------

...

1086	0	0:00	/sbin/busybox ps fauxwww
------	---	------	--------------------------

# Debugger

---

- The main process handles communication
- SDK functions (using standard communication channels) can't be used
- *Auxiliary* channels available (eg. shared memory)
- gdb on the host, gdbserver on the HSM
- Additional challenge:
  - The main process is monitored with *ptrace*
  - Reboots the HSM in case of crashes

## Information gathered

---

- Dynamic analysis ability
  - Every process run as *root*
  - No hardening nor mitigation options
- The bootloader is a slightly modified version of U-Boot:
  - No secure boot mechanism


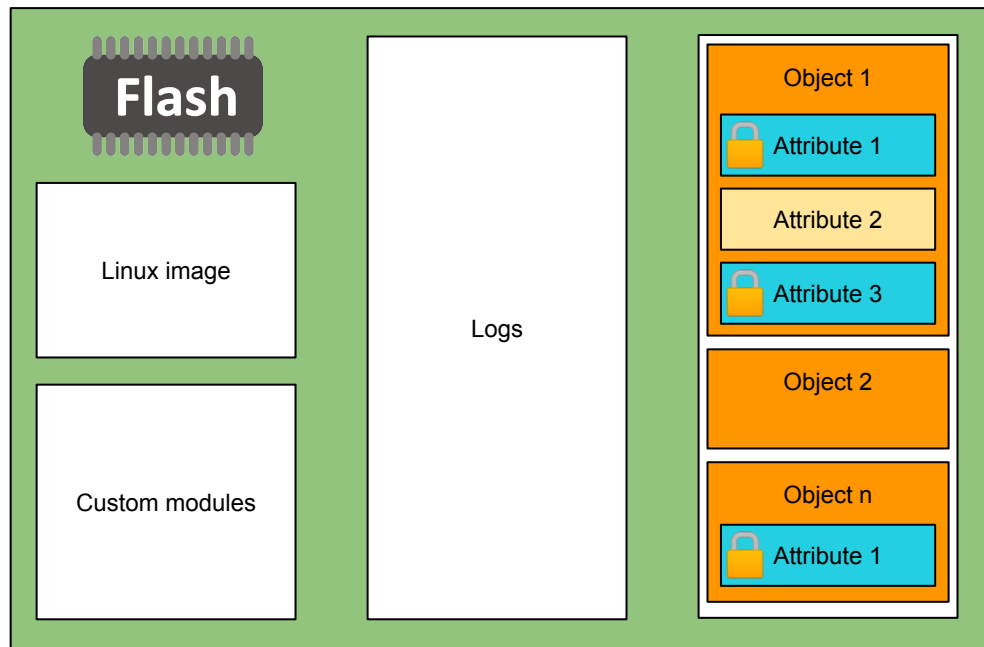
## Storage

---

- Persistent data stored to a 64 Mb flash memory on the PCI card:
  - Linux image
  - Custom modules
  - Logs
  - PKCS #11 objects
- PKCS #11 objects and authentication information are stored into a dedicated partition using a proprietary filesystem
- Sensitive object attributes are stored encrypted and decrypted on-the-fly

# Storage

---



## Storage

---

- No logical separation across HSM slots
  - Each objects from each slots are stored in the same flash partition
  - Reverse engineering shows that secrets are stored with the same key
- Code execution on the HSM allows to dump all secrets



# First Code Execution

---

## Grepping for memcpy

---

- ~700 calls to memcpy in API .so
- Manual analysis:
  - memcpy called from PKCS #11 functions
  - Variable size parameter and stack destination
- MilenageDerive is the only one vulnerable to a stack overflow
- CKM\_MILENAGE\_DERIVE mechanism:
  - UMTS (Universal Mobile Telecommunication System) authentication algorithms
  - Used by HSMs in Telco environment
  - Key derivation for  $f_3$ ,  $f_4$ ,  $f_5$  and  $f_5^*$  MILENAGE functions

## Bug

---

- CKM\_MILENAGE\_DERIVE requires a handle on a 16-byte MILENAGE key
- Stored as a generic secret key, installed with CreateSecretKey
- MilenageDerive does not check the length of the secret key

```
int MilenageDerive(...) {
    uint8_t aesKey[16];
    ...
    GetObjectClassAndKeyType(keyObject, &attributeClass, &keyType);
    if (attributeClass == CKO_SECRET_KEY) {
        keyValue = FindAttr(CKA_VALUE, keyObject);
        if (keyValue) {
            valueLen = keyValue->valueLen;
            memcpy(aesKey, keyValue->pValue, keyValue->valueLen);
            ...
        }
    }
```


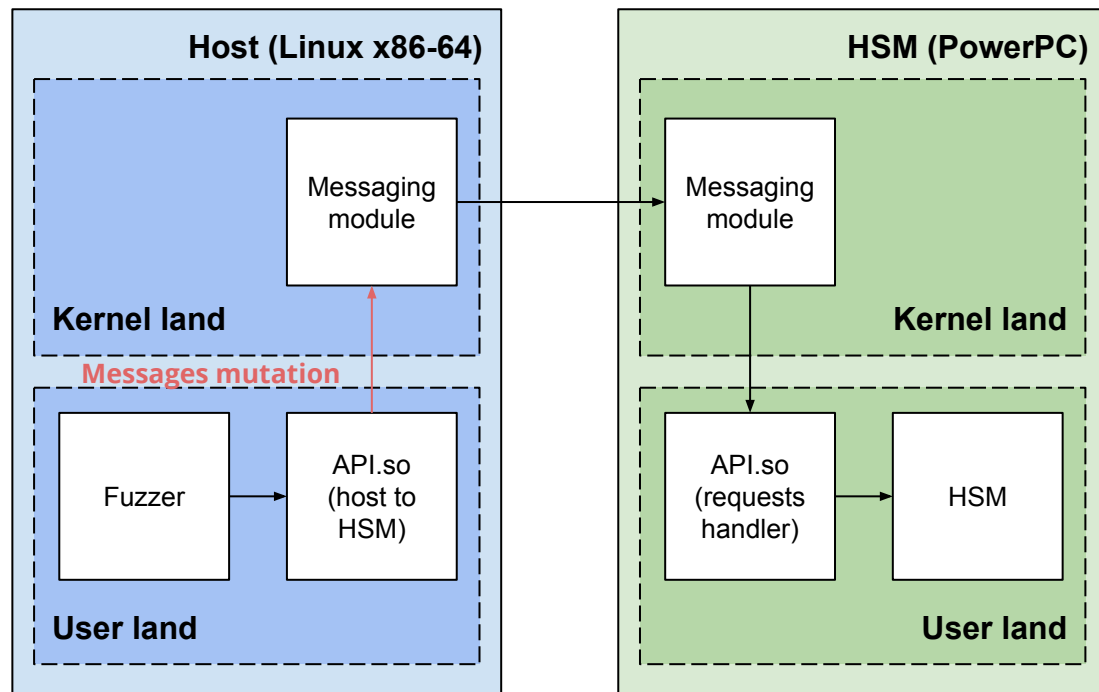
# Exploit

---

- Code exec is trivial
  - No stack cookie
  - No ASLR
- But
  - Resuming execution is tricky
  - `CreateSecretKey` requires to be authenticated
  - MILENAGE is present only on recent firmware versions

→ Better look for another bug

# Fuzzing



# Fuzzing

---

- Testsuite from PKCS #11 usage examples
- Random bytes mutation
- Main challenges:
  - Host kernel module crashes
  - HSM Denial of Service due to OOM
- Results:
  - 14 vulnerabilities, several classes of memory corruption bugs
  - *Heartbleed*-like vulnerability
  - Stack and heap overflows
  - etc.

# Heartbleed

- Memory leak of the HSM's heap
- Authentication required

```
attacker@host:~$ ./heartbleed user $((0x78)) | hd
[*] modifying buffer size to 0x78
00000000  62 6c 61 68 00 90 47 6c |blah..Gl|
00000008  00 00 00 04 01 00 00 00 |.....|
00000010  01 00 00 00 01 00 00 00 |.....|
00000018  01 01 01 00 00 01 01 00 |.....|
00000020  00 00 08 01 73 75 62 6a |....subj|
00000028  65 63 74 00 00 00 01 02 |ect.....|
00000030  00 00 00 01 01 00 00 00 |.....|
00000038  00 11 00 00 00 08 01 10 |.....|
00000040  43 d8 5c 37 a7 57 6b 00 |C.\7.Wk.|
00000048  00 01 61 00 00 00 04 01 |..a.....|
00000050  00 00 00 01 80 00 01 02 |.....|
00000058  00 00 00 10 01 32 30 31 |.....201|
00000060  38 30 39 30 33 30 38 30 |80903080|
00000068  33 34 39 30 30 00 00 01 |34900...|
00000070  0a 00 00 00 01 01 01 80 |.....|
```

# Reliable Code Execution

---



## Bug Discovery

---

```
CK_MECHANISM digestMechanism = { CKM_RIPEMD128, NULL_PTR, 0 };
unsigned char state[4096], data[32];
CK_ULONG ulStateLen;

C_DigestInit(hSession, &digestMechanism);
C_GetOperationState(hSession, state, &ulStateLen);
mutate(state, ulStateLen);
C_SetOperationState(hSession, state, ulStateLen, 0, 0);
C_DigestUpdate(hSession, data, sizeof(data));
```

# Crash Analysis

- A single byte mutation triggers a crash during restore
- NULL-deref but unusual stacktrace
- Static and dynamic analysis
- Type confusion bug:
  - The mutated byte describes the object type
  - An unexpected digest object can be restored
  - Object A's methods can be called object B

[illegible]

# Exploit Development

---

- Digest mechanisms analysis:
  - Memory leak
  - Relative write primitive
- Complex but reliable exploit
  - Heap feng shui
  - Shellcode across various and not consecutive objects
  - Cache coherency
  - etc.

## Payload issue

---

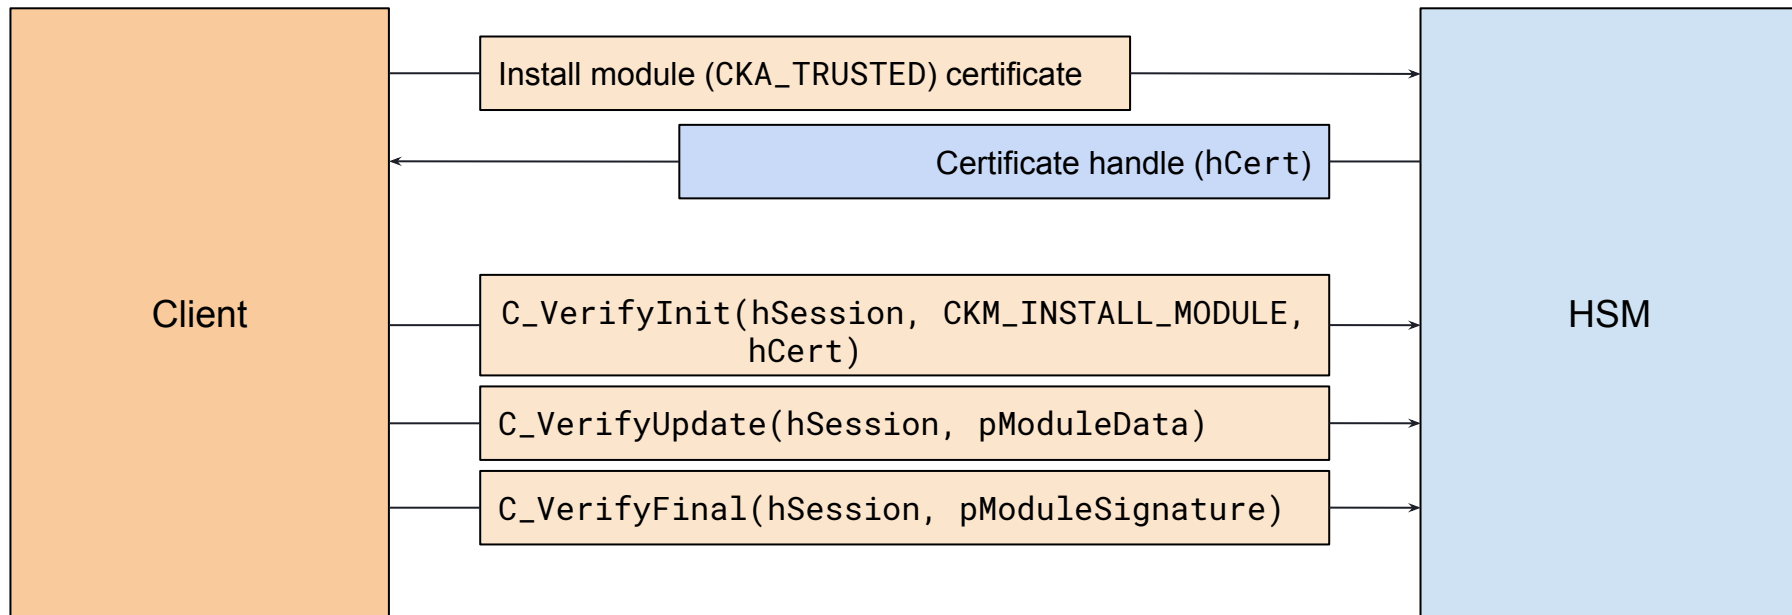
- Payload executed with root privileges
- `system("/bin/sh")` shellcodes won't work
  - No interesting binary on the HSM to execute
  - No simple way to communicate between the host and the HSM
- Final payload
  - a. Patch of the PIN verification function
  - b. Login as admin
  - c. Evil module installation: dump of the flash and the decryption key
  - d. Offline decryption
- The exploit is a single binary executed from the host

# Firmware Signature Bypass

---

# Module Install

---



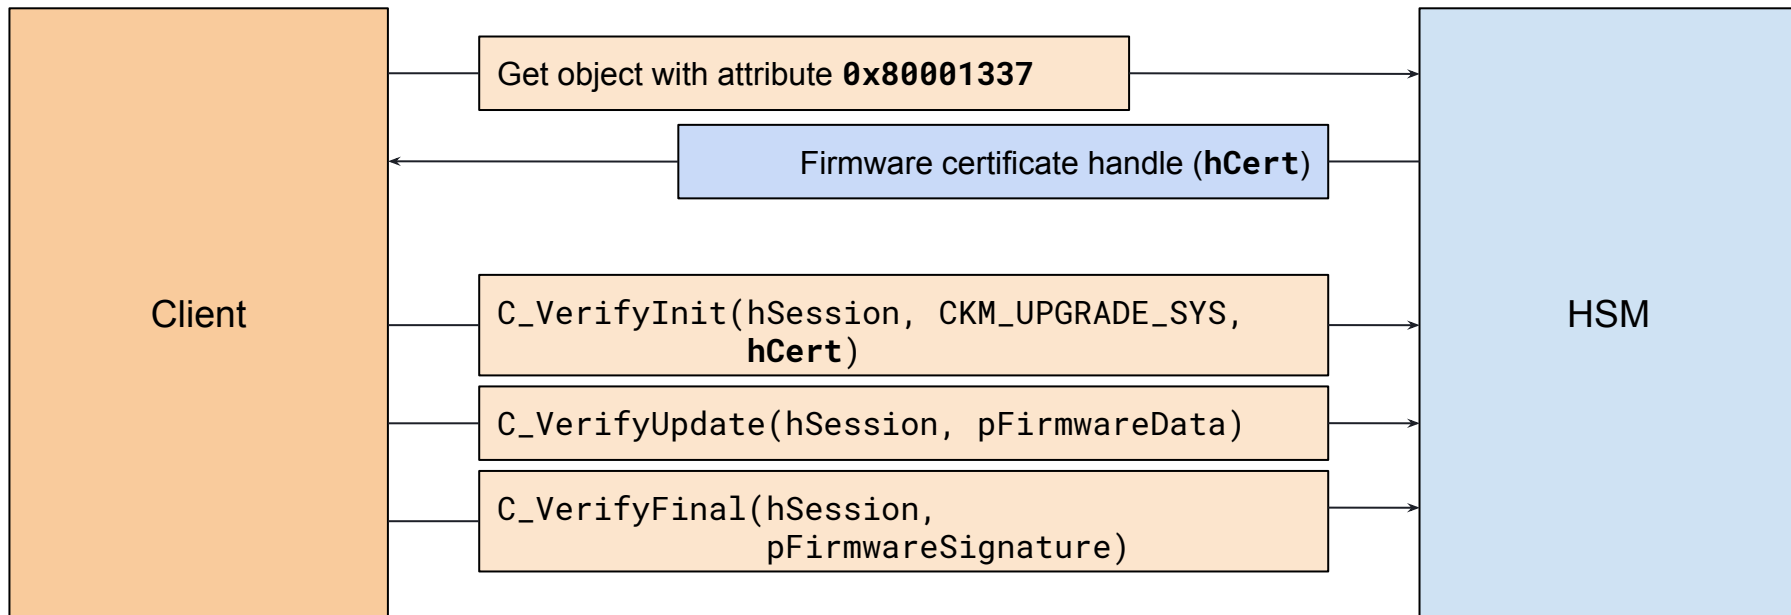
# Firmware Update

---

- Almost identical to modules install
- Firmware updates are signed by the vendor
  - Ensures integrity
  - No way to install a custom firmware
- Vendor certificate hardcoded in the (installed) firmware code
- `admin@hsm-host:~$ vendor-fw-update /tmp/firm-1.3.bin`

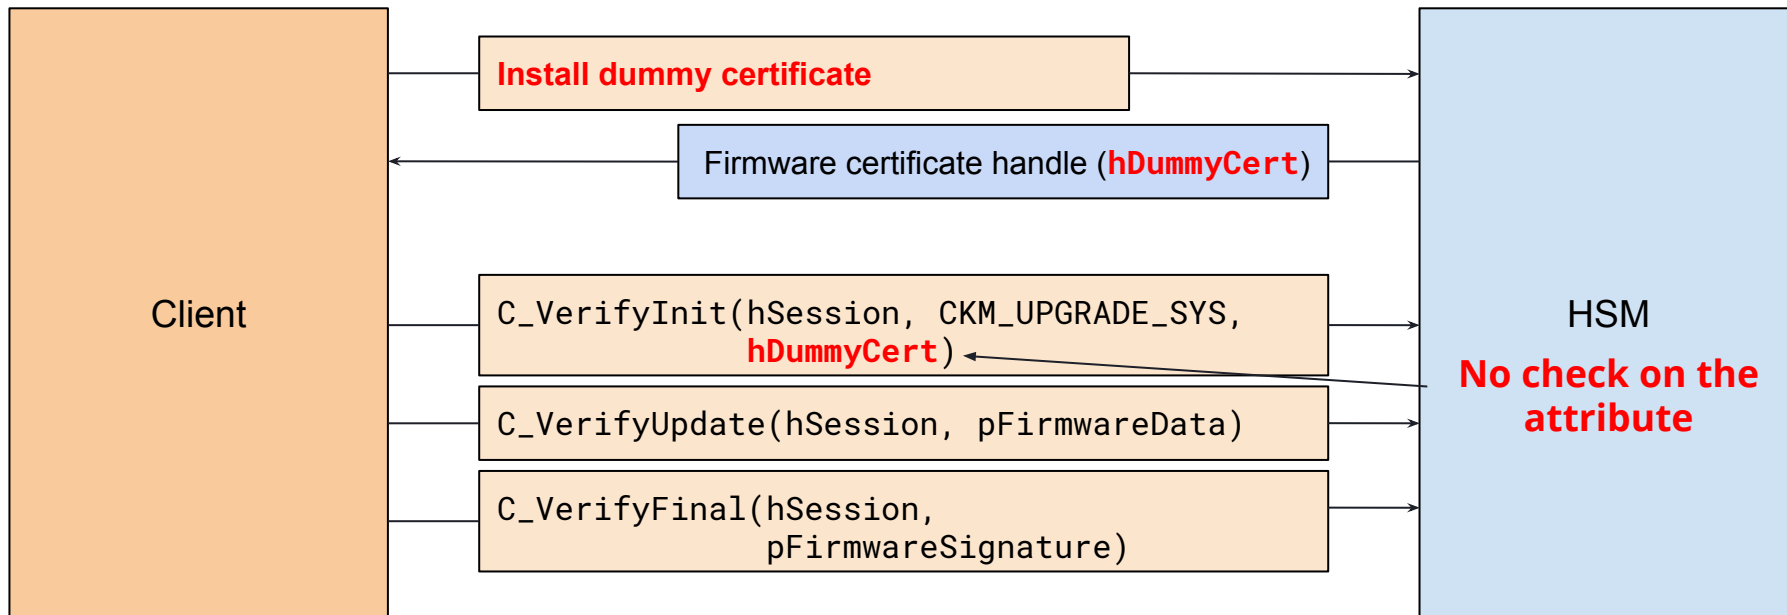
# Firmware Install

---





## Firmware Install without vendor signature



## Result

---

- Malicious firmware update
- Persistent backdoor
- Downgrade firmware attack
- (Requires admin privileges)

# Conclusion

---

## Arbitrary Code Execution

---

- Vulnerability research against unauthenticated PKCS #11 operations
- Several memory corruption vulnerabilities found
- Pre-auth reliable exploit
- Consequence: arbitrary code execution on the HSM
- Does it work against net HSMs?

## Secret Decryption

---

- Dump of the storage (containing encrypted secrets)
- Dump of the encryption key
- Offline decryption of the HSM secrets

# Persistence

---

- Malicious firmware installation using either:
  - The signature bypass
  - Code execution
- The HSM integrity cannot be guaranteed anymore:
  - No secure boot
  - This vulnerability can be exploited again because of downgrade attacks

## Responsible disclosure

---

- Every vulnerability reported to the vendor
- New firmware update
- Pay attention to your vendor security advisories and apply updates

## Key Takeaways

---

- Not an exhaustive HSM study: what about other models and other vendors?
- Methodology to look for vulnerabilities, improve the overall security of the industry
- HSMs mostly certified against hardware attacks, what about software?





Questions?

