

Binary symbolic execution with KLEE-Native

Sai Vegasena

The Team



Sai Vegasena

Security Engineering Intern

sai.vegasena@trailofbits.com
@svegas18



Peter Goodman

Senior Security Engineer

peter@trailofbits.com
@peter_a_goodman



During my internship I ...

- Developed a fork of KLEE that operates on raw binaries
- Translated machine code to LLVM bitcode with Remill and ran it in KLEE
- Wrote a custom allocator to accurately model heap memory in KLEE's emulator
- Implemented virtualized system calls that productively handled symbolic data
- Developed a new forking model for KLEE's symbolic executor
- Reproduced a CVE used in an old ChromeOS exploit chain

KLEE is a symbolic virtual machine that executes LLVM bitcode



- Applied in software testing and verification
- Dynamically generates high-coverage producing inputs
- Leverages a custom runtime environment



Bugs Galore :)

KLEE's greatest strength is also its greatest weakness


Pros of running LLVM bitcode

- Allows for custom runtime definitions and intrinsics
- Executes anything clang can compile
 - i.e C, C++, Rust, Swift, Go,, etc

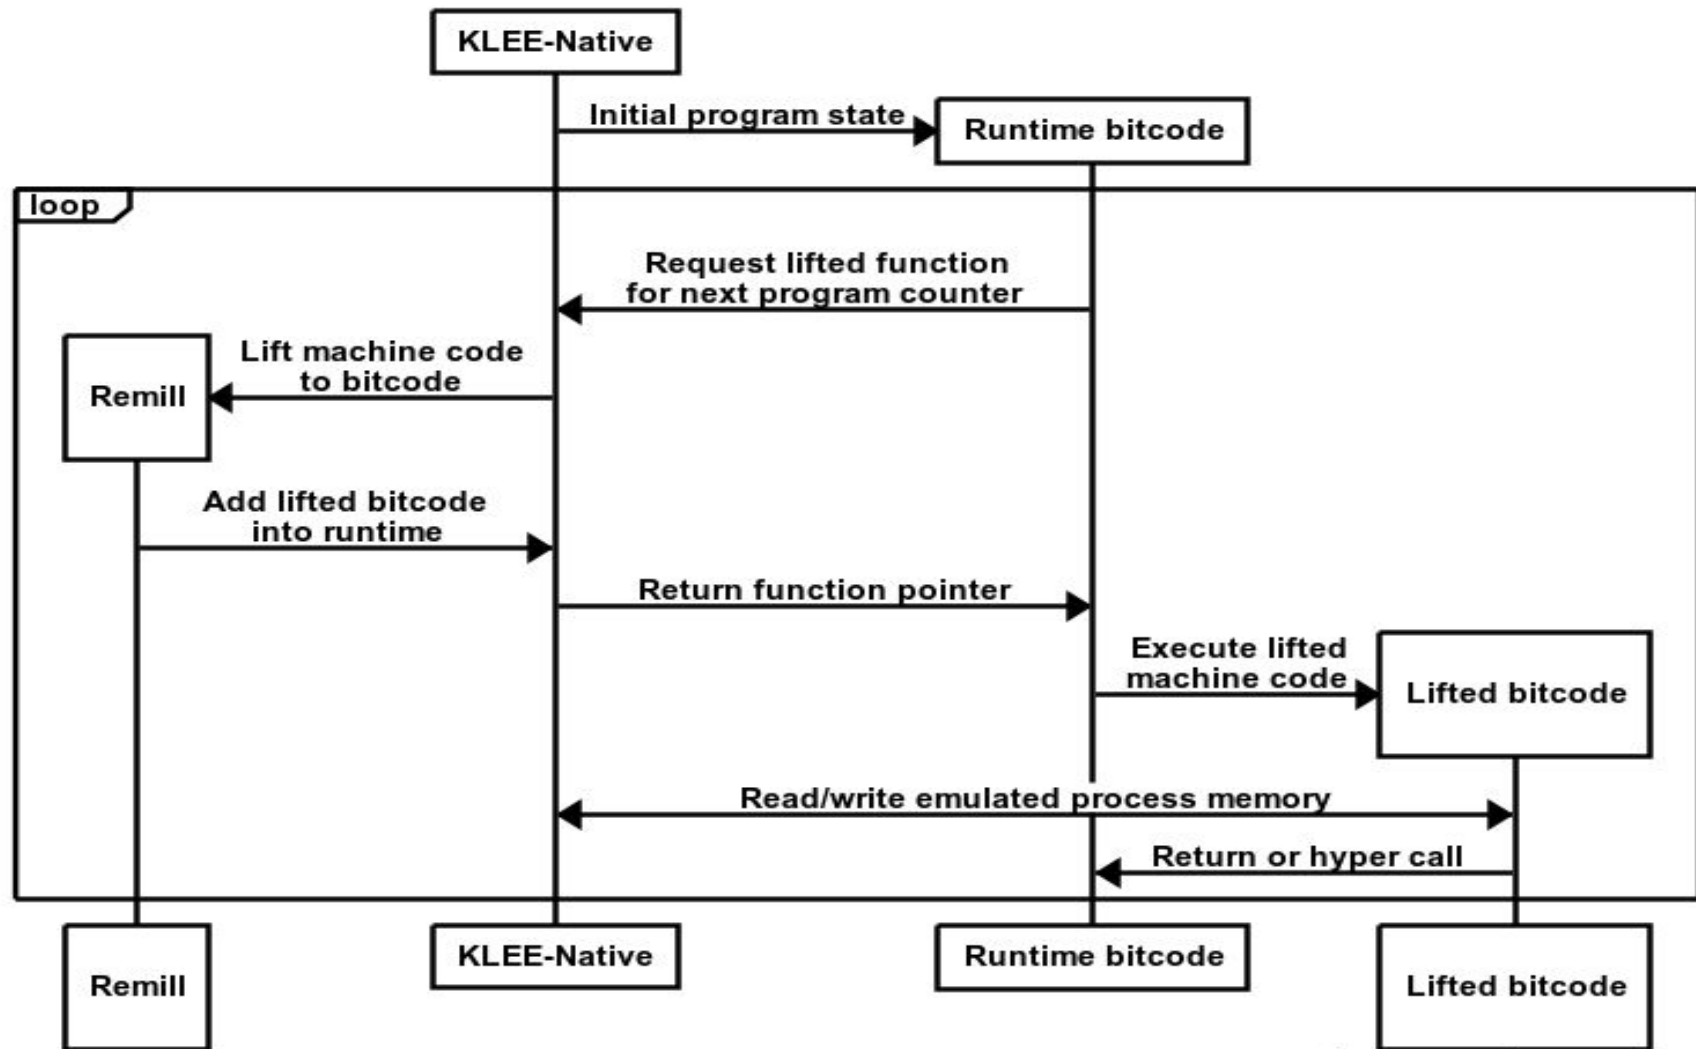
Cons of running LLVM bitcode

- Sometimes need source
- Build systems, configs, and dependencies
- Manually Injecting KLEE-API calls into the source
- McSema is an option but CFG recovery is limiting and there are occasional inaccuracies

KLEE-Native operates on snapshotted program binaries

- Binaries are snapshotted with user defined breakpoints
 - Static breakpoints
 - Dynamic breakpoints for ASLR
 - `./klee-snapshot-7.0 --workspace_dir ws --dynamic --breakpoint 0x1337 --arch amd64_avx -- ./a.out`
-  lifts machine code instructions to LLVM bitcode
- KLEE-Native executes the runtime and the lifted LLVM

`./klee-exec-7.0 --workspace_dir ws`



Runtime is the kernel and the machine

- Remill async hyper call is defined in runtime
 - “Implements” OS functionality
 - Execution is passed to a linux system call wrapper in runtime
- Custom ABI extracts arch info from state
 - Store return value
 - Extract args
 - Find syscall number
- Wrappers do error checking done by OS
 - Gives a kernel-level “insight”

```
template <typename ABI>
static Memory *SysOpen(Memory *memory, State *state,
                       const ABI &syscall);
```

```
if (path_len >= PATH_MAX) {
    STRACE_ERROR(open, "Path name too long: %s", gPath);
    return syscall.SetReturn(memory, state, -ENAMETOOLONG);
    // The string read does not end in a NUL-terminator; i.e. we read less
    // than 'PATH_MAX', but as much as we could without faulting, and we didn't
    // read the NUL char.
} else if ('\0' != gPath[path_len]) {
    STRACE_ERROR(open, "Non-NUL-terminated path");
    return syscall.SetReturn(memory, state, -EFAULT);
}
```


Simplifying lifted libc functions with accuracy

- Problem
 - Lifting libc functions is slow
 - Unnecessary state forking on symbolic data
- LD_PRELOAD-based library into snapshotted programs
 - Lets us interpose and hook to simple libc variant in the runtime
- Variants handle symbolic data in a simple way with no lifting

Heap memory is accurately modeled with libc intercepts

- Lifted mallocs call brk or mmap
 - Technically accurate but bad for bug-finding
 - No clarity for bounds checks on allocations
 - Hard to oversee UAFs, double frees, and access violations
- Utilize the intercept hook to organize allocations in a uniform structure
- Basically implemented a custom allocator
- Custom address encoding helps “locate” allocations in alloc lists on mallocs and frees

Encoding

```
union Address {  
    uint64_t flat;  
    struct {  
        uint64_t offset :16;  
        uint64_t must_be_0x1 :4;  
        uint64_t size :16;  
        uint64_t alloc_index :24;  
        uint64_t must_be_0xa :4;  
    } __attribute__((packed));  
};
```

It is possible to fall back to real libc functions in KLEE-Native

Before Snapshot

→ `malloc:`

```

; CODE XREF: .text:
; .text:__GI___libc

```

Emulated in KLEE-Native

→ `handle_malloc_hypercall:`

```

    jmp     cs:ptr_to_malloc_handler
    int     81h
    retn

```

Fallback to real libc malloc just in case

→ `handle_malloc_passthrough:`

```

    push    r10
    mov     r10, cs:ptr_to_addr_of_real_malloc
    mov     r10, [r10]
    xchg    r10, [rsp]
    retn

```

It is possible to fall back to real libc functions in KLEE-Native

Before Snapshot

→ `malloc:`

```
; CODE XREF: .text:
: .text: GT __libc
```

Emulated in KLEE-Native

NOP when we load state

→ `handle_malloc_hypercall:`
`int 81h`
`retn`

Fallback to real libc malloc just in case

→ `handle_malloc_passthrough:`
`push r10`
`mov r10, cs:ptr_to_addr_of_real_malloc`
`mov r10, [r10]`
`xchg r10, [rsp]`
`retn`

It is possible to fall back to real libc functions in KLEE-Native

Before Snapshot

→ `malloc:`

```
; CODE XREF: .text:
: .text: GT __libc
```

Emulated in KLEE-Native

→ `handle_malloc_hypercall:`

NOP when we load state

```
int
retn 81h
```

Fallback to real libc malloc just in case

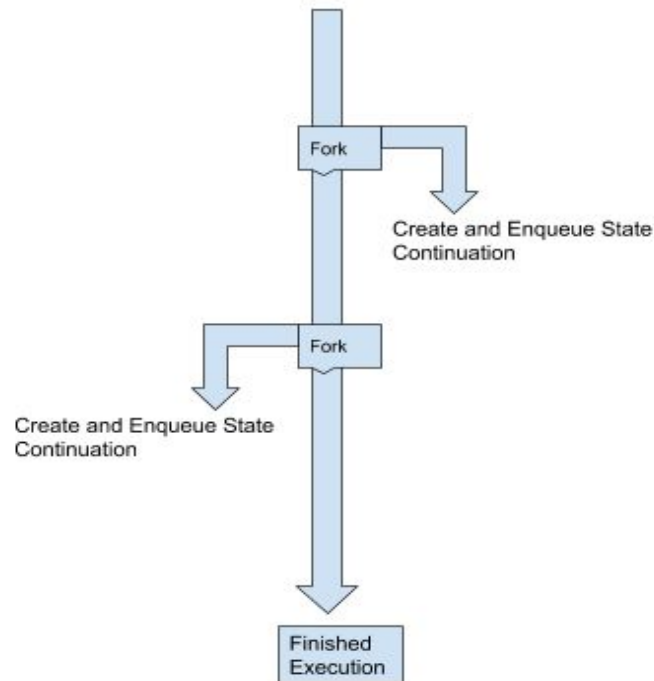
→ `handle_malloc_passthrough:`

```
push    r10
mov     r10, cs:ptr_to_addr_of_real_malloc
mov     r10, [r10]
xchg    r10, [rsp]
retn
```

Skip the ret in emulator

Eager concretization is better than eager forking

- Closer in spirit to SAGE, a static symbolic symbolic executor
- Akin to lazy evaluation
- “State continuations” are moral equivalent of a Python generator, and may be invoked to give “the next viable fork at this point”
- Continuations are enqueued and scheduled later
- Handle branches and symbolic addresses



KLEE-Native can identify real bugs

- Make a snapshot in vulnerable function

```
klee-snapshot-7.0 --workspace_dir ws_CVE --dynamic --breakpoint  
0xb33 --arch amd64_avx -- ./c_ares_repro
```

- Run klee-exec

```
klee-exec-7.0 --workspace_dir ws_CVE
```

- Policy handler provides an interesting feature that Valgrind and Asan don't


```

KLEE: WARNING ONCE: Alignment of memory from call "_Znwm" is not modelled. Using alignment of 8.
KLEE: WARNING ONCE: calling external: vfprintf(140017979533152, 94410329517216, 94410333418800) at [no debug info]
libc_memcpy:dest=7ffcee613f70, src=400d94, len=8, ret=7ffcee613f70
libc_malloc:size=19, ptr=a000000001310000
E0729 16:48:17.126665 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9dfd9d 3 (BYTES 0f 50 dc) MO
VMSKPS_GPR32_XMMps (WRITE_OP (REG_64 RBX)) (READ_OP (REG_128 XMM4)))
E0729 16:48:17.128990 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9e01c2 4 (BYTES 66 0f 50 d8)
MOVMSKPD_GPR32_XMMpd (WRITE_OP (REG_64 RBX)) (READ_OP (REG_128 XMM0)))
E0729 16:48:17.130098 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9e0434 3 (BYTES 0f 50 dd) MO
VMSKPS_GPR32_XMMps (WRITE_OP (REG_64 RBX)) (READ_OP (REG_128 XMM5)))
E0729 16:48:17.140801 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9e1a1a 4 (BYTES 66 0f 50 d8)
MOVMSKPD_GPR32_XMMpd (WRITE_OP (REG_64 RBX)) (READ_OP (REG_128 XMM0)))
E0729 16:48:17.174049 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9dfd9d 3 (BYTES 0f 50 dc) MO
VMSKPS_GPR32_XMMps (WRITE_OP (REG_64 RBX)) (READ_OP (REG_128 XMM4)))
E0729 16:48:17.176441 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9e01c2 4 (BYTES 66 0f 50 d8)
MOVMSKPD_GPR32_XMMpd (WRITE_OP (REG_64 RBX)) (READ_OP (REG_128 XMM0)))
E0729 16:48:17.177579 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9e0434 3 (BYTES 0f 50 dd) MO
VMSKPS_GPR32_XMMps (WRITE_OP (REG_64 RBX)) (READ_OP (REG_128 XMM5)))
E0729 16:48:17.188436 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9e1a1a 4 (BYTES 66 0f 50 d8)
MOVMSKPD_GPR32_XMMpd (WRITE_OP (REG_64 RBX)) (READ_OP (REG_128 XMM0)))
E0729 16:48:17.305361 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb9f6250 8 (BYTES 66 0f 57 05
b8 6b 14 00) XORPD_XMMxuxq_MEMxuxq (WRITE_OP (REG_512 ZMM0)) (READ_OP (REG_512 ZMM0)) (READ_OP (DWORD_PTR (ADD (REG_64 PC
) (SIGNED_IMM_64 0x146bc0))))))
E0729 16:48:17.633155 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb1abca 4 (BYTES d9 74 24 d8)
FNSTENV_MEMmem28 (WRITE_OP (DWORD_PTR (ADD (REG_64 RSP) (SIGNED_IMM_64 -0x28))))))
E0729 16:48:17.633316 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb1abfb 4 (BYTES d9 74 24 d8)
FNSTENV_MEMmem28 (WRITE_OP (DWORD_PTR (ADD (REG_64 RSP) (SIGNED_IMM_64 -0x28))))))
E0729 16:48:17.633388 6124 Lifter.cpp:123] Missing semantics for instruction (AMD64 7fbfcb1ac10 4 (BYTES d9 74 24 d8)
FNSTENV_MEMmem28 (WRITE_OP (DWORD_PTR (ADD (REG_64 RSP) (SIGNED_IMM_64 -0x28))))))
KLEE: WARNING ONCE: Alignment of memory from call "_Znam" is not modelled. Using alignment of 8.
KLEE: WARNING ONCE: calling external: write(2, 94410361759936, 12) at [no debug info]
HIT STRCMP!
1:write:fd=2, size=12/12
E0729 16:49:12.509608 6124 AllocList.cpp:157] Heap address overflow on memory write address a000000001310013
KLEE: ERROR: (location information missing) Failed 1-byte write of 1 to address 0xa000000001310013 in address space 1
KLEE: NOTE: now ignoring this error at this location
I0729 16:49:12.509789 6124 Executor.cpp:3078] Finished state, continuation stack size is 1

```

Sai Vegasena: sai.vegasena@trailofbits.com

Peter Goodman: peter@trailofbits.com

Symex on Binary Snapshots with KLEE-Native

