

# Getting started with osquery

Lauren Pearl & Andy Ying  
Engineering team @ Trail of Bits

# Agenda




- Trail of Bits
- Intro to osquery
- How to contribute to osquery
- Q&A

**Cyber security research company** - High-end security research with a real--world attacker mentality to reduce risk and fortify code.


## Security Research

- As a leading cybersecurity research provider to DARPA, the Army and the Navy – we create and release open source research tools




## Security Engineering

- We offer custom engineering for every stage of software creation, from initial planning to enhancing the security of completed works




## Security Assessments

- We offer security auditing for code and systems requiring extreme robustness and niche system expertise



## Security Engineering

- We offer custom engineering for every stage of software creation, from initial planning to enhancing the security of completed works




Open source engineering examples:



Google Santa




StreamAlert



# osquery: The basics

TRAIL  
OF  
BITS

# What is osquery?



Open source endpoint detection tool, developed by **Facebook**, that allows users to query their fleets' system data like they would a relational database.

# osquery - What can you ask it?

```
osquery> SELECT * FROM mounts m, disk_encryption d
WHERE m.device_alias = d.name
AND m.path = "/"
AND d.encrypted = 0;

      device = /dev/disk1
device_alias = /dev/disk1
```

## Primary disks that are unencrypted

Disk encryption is an organizational policy for many companies, this query returns hosts whose primary disk is currently unencrypted.

[See Available Tables](#)

## Other query examples:

- Did anyone download a file called "XYZ" today?
- Tell me all the users that have accessed root
- What processes are running on disk without a binary?

## How queries run:

- **Osqueryi:**
  - One-off
- **Osqueryd (daemon):**
  - Scheduled (query once per hour)
  - Event-triggered (FIM, Process Auditing)

# How to write queries



Format:

```
> SELECT columns FROM table  
WHERE parameters;
```

Rules of  
the Road:

- Don't use SELECT \* without parameters for a data-heavy table!
  - ~~SELECT \* FROM processes;~~
- Read up on [SQL constraints](#)!
- Start from queries in [QueryPacks](#), [documentation](#), and blogs!
- Test queries on your local machine before running on your fleet!



# Tables & schemas

Currently 226 tables to join!

- Each table has a different schema
- You can view all available tables and their schemas on the [osquery website](https://osquery.org)
- Or you can find either using in-line commands:

`.schema [TABLE]` **Show the CREATE statements**

`.tables [TABLE]` **List names of tables**

Some popular examples:

```
•acpi_tables      •processes      •rpm_packages
•arp_cache        •routes         •apt_sources
•crontab          •shell_history  •deb_packages
•file_events      •smbios_tables  •homebrew_packages
•kernel_info      •suid_bin       •kernel_modules
•listening_ports  •system_controls •memory_map
•logged_in_users  •usb_devices    •shared_memory
•mounts           •users          •browser_plugins
•pci_devices      •groups         •startup_items
```

# Live Example:

Hey osquery, is my machine running hot?

```
SELECT * FROM temperature_sensors;
```

# Contributing

TRAIL  
OF BITS

# Extending osquery - Core or extensions?

Does my contribution belong in Core or in an Extension?

## **Belongs in Core:**

- Observes guiding principles
- Has been shared with and approved by osquery project maintainers as a new feature in Core
- Meets Facebook's testing and quality standards

## **Belongs in an extension:**

- Might not observe the osquery core guiding principles
- Has not been shared with or approved by Facebook as a new feature in Core
- Expands the scope of use for osquery beyond endpoint monitoring
- Integrates with a proprietary or esoteric tool that is not widely applicable

# Contributing to Core

## Guidelines for contributing features to osquery core

While there are occasional exceptions, contributions to core should abide by the following osquery guiding principles in order to be accepted:

1. osquery doesn't change the state of the system
2. osquery doesn't create network traffic to third parties
3. osquery's endpoint binaries have a light memory footprint
4. osquery minimizes system overhead & maximizes performance
5. The query schema for osquery seeks uniformity between operating systems





**Freeze!**

**Osquery core development will be frozen  
until 11/19/2018**



# Contributing to Core



1. Read the [CONTRIBUTING.md](#) doc on osquery github repo!
2. Leverage the [issue submission templates](#)
3. For significant changes start by opening a [Blueprint](#) issue which describes your idea, the problem you're solving and how you plan to implement your solution
4. Single PRs should represent a single body of work - never several unrelated changes (this includes required bug fixes)
5. Start by developing on a feature branch, then submit a pull request against the osquery master branch
6. Tag any related issues in the PR description
7. During testing, submit updated code to comments rather than editing initial PR

# Contributing via extensions

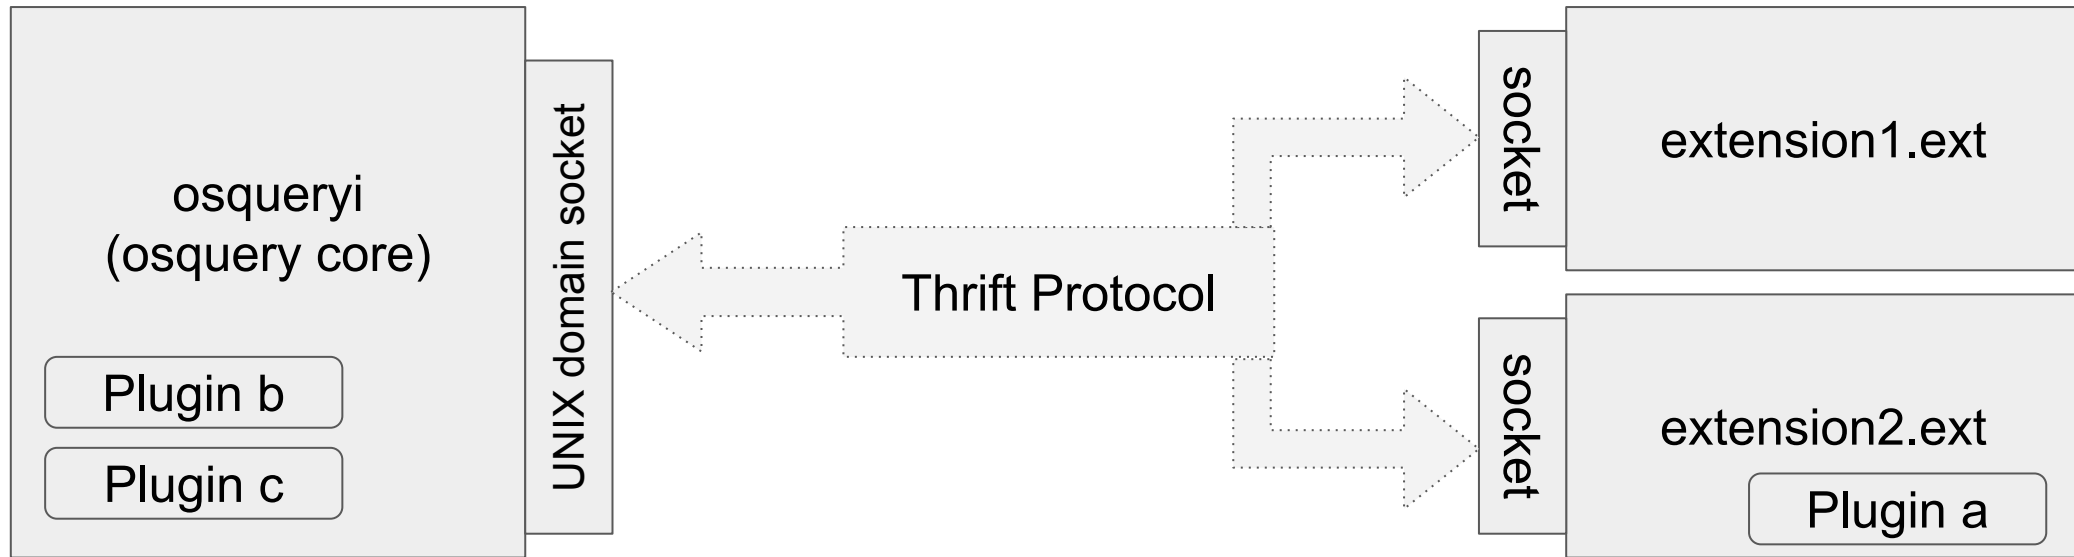


**Extensions** are the way to develop custom functionality in a separate component that *extends* (via virtual tables) or *overrides* (via plugins) the osquery core behavior.

Extensions compile and run as separate executables. They communicate with the osquery core process using the **Thrift** RPC protocol.



# osquery extensions



```
osqueryi --extension /path/to/extension1.ext --extension /path/to/extension2.ext
```

# Building an osquery extension (C++)

1. Statically link the core code
2. `#include <osquery/sdk.h>`
3. Inherit from `TablePlugin` and implement appropriate methods
4. Register the table or plugin
5. Initialize osquery worker threads
6. `startExtension()` to connect to osquery core
7. Symlink your build directory into osquery's
8. `make externals` or  
`cmake --build . --config Release --target external_EXTENSION_NAME`

# Building an osquery extension (Python)

<https://github.com/osquery/osquery-python>

1. `> pip install osquery`
2. `import osquery`
3. `@osquery.register_plugin`
4. `class MyTablePlugin(osquery.TablePlugin)`
5. `def name(), def columns(), def generate()`
6. `osquery.start_extension(name="my_ex", version="1.0.0")`

```
osqueryi --extension path_to_my_table_plugin.py
```

# Some useful osquery resources

- **Super-simple guide for getting started:** Scott Roberts - *osquery 101 — Getting Started*  
<https://medium.com/@sroberts/osquery-101-getting-started-78e063c4e2f7>
- **Use cases for the security team:** Chris Long - *osquery For Security*  
<https://medium.com/@clong/osquery-for-security-b66fffd2daf>
- **How to build extensions in C++ for Windows, straight from the source:** Nick Anderson - *Building and deploying osquery extensions on Windows*  
<https://brewfault.io/blog/2018/1/29/building-extensions-for-osquery-on-windows>
- **The future - what osquery users want next:** Trail of Bits - *What do you wish osquery could do?*  
<https://blog.trailofbits.com/2018/04/10/what-do-you-wish-osquery-could-do/>

# Contact Us




**Lauren Pearl**

Head of Strategy and Ops

---

[lauren@trailofbits.com](mailto:lauren@trailofbits.com)



**Andy Ying**

Security Engineer

---

[andy@trailofbits.com](mailto:andy@trailofbits.com)