

TRAIL *OF* BITS

Property-based testing of smart contracts

Solidity testing

- “Is our code correct?”
- Typically:
 - Specify some sequence of interactions
 - Check the results
 - Run this every commit/release
 - “First, call **approve**, then **transferFrom**, then **balances** should...”

When doesn't this work?

- Problem: we don't test most behavior
- “[...] approve, then transferFrom, then [...]”
 - What about in the other order?
 - What about calling **approve** twice?
 - Writing tests for everything is hard

Can we write better tests?

- Tests so far specify a single input
 - `f(3) == 1337`
 - `[1, 4, 2, 13].sort() == [1, 13, 2, 4]`
- ***Property tests* cover any possible input**
 - $\forall x: \text{leet}(f(x)) \geq \text{leet}(x)$
 - $\forall l: l.\text{sort}()$ must be alphabetical

What's the difference?

- Unit tests cover cases devs know about
- Property tests cover cases they don't
 - Easy way to get test coverage up
 - Finds weeiird edge cases
 - History of being unreasonably effective

OK, now do smart contracts

- **Unit tests specify a single transaction sequence**
 - If alice calls f, sends bob 2 ETH, then calls g, she'll have 3 ETH
- **Property tests cover any possible transaction sequence**
 - No matter what methods alice calls, bob can't lose money

How do we test like this?

- The set of possible inputs is *giant*
- Either we reason about some of them, or all of them
 - Reason about some of them: Echidna
 - Reason about all of them: Manticore
- **One test, two ways to check**

- **Given an ABI, generates random transactions**
 - function `f(uint x, uint y) [...]`
 - generate 10,000 pairs of uints
 - call `f` with all of them
 - check the property on each
- **No *guarantee*, better than one sequence**

- Implements a superset of EVM
- Regular EVM: values can be a number
 - Examples: 0, $2^{256} - 1$, "hello world" (encoded)
- **Manticore: a number, or "all numbers such that..."**
 - Examples: "anything > 3 , < 17 ", "any prime number", "literally anything"
 - Represented as constraints, so z3 can solve
- **Manticore : Geth :: Property testing : unit testing**

- Execute with “any possible initial transaction”
- See if failure is ever possible
- If not, your property holds, congrats!
- If yes, use Z3 to solve constraints

What's the difference?

- **Manticore: complicated, but comprehensive**
 - Lots of effort to set up
 - Effectively formal verification
 - Super high assurance
- **Echidna: pretty easy, but it's random**
 - Runs more or less automatically
 - You could miss stuff though

How should I property test?

- Fun personal projects: don't worry about it
- No existing tests: write some unit tests
- Existing unit tests: use Echidna
- Good test suite/*needs* correctness/\$\$\$: use Manticore

Questions?



JP Smith

Security Engineer

jp@trailofbits.com, @japesinator

www.trailofbits.com